

SWC DB (Super Wide Column Database)

The database is structured in columns, not by Tables or Namespaces, as familiar SQL is 'select columns from "table_name";', with "Super Wide Column DB" It is 'select [where_clause [Columns-Intervals]]; '.

The "Super Wide Column" comes to its meaning by the column's key is a list-set of keys eg. A column cell's key is keys=(k1,k2,k3,k4,kN), comparing to a "Wide Column" key that is row, column-family, column-family-qualifier in the "Super Wide Column DB" it is equal to keys=(k1(row),k2(cf),k3(cq)) or(similar) column=cf with keys=(k1(row),k3(cq)) . Majority of the developments are planned to be on bases of Hypertable (<https://github.com/kashirin-alex/hypertable>).

The storage-form in the "Super Wide Column DB" is based on column-id and range-id, which on path consist CellStores and CommitLogs files at any point one server is responsible for a range-id on column-id.

The CellStores are Files storing Cells in serialized form that are after latest compaction whereas CommitLogs are the open-file-descriptor to which current data is added.

The Serialization of data in a CellStore/CommitLog file: (delimited with "|" for visual-representation of NONE)

|Blocks(Header | Compressor(Cells)) | Fixed-Index | Variable-Index | CellStore-Trailer|

The Cell-Serialization: |Key-length(int32)|Key-serialized| Value-length(int32)|Value-Data|

Key serialization: |Key-flag(int8)|Key-control(int8)|Keys-count(int8)|joined(Keys[N]\0)|Timestamp(int64)|Revision(int64)|

The Cells Ranges, a range is a Keys-start to Keys-end, SWC DB use a self-explanatory master-ranges that define ranges to meta-ranges of data-ranges(cells-range) whereas on range-locator it includes the Keys comparison on the comparators of request, resulting in as most narrowed scan of cells-ranges.

System's reserved columns id[1-9],

1: IDENTIFIER a counter type column

2: RANGES

3: RID(RANGE-ID) to RS(N)

4: Column ID, Column Name and serialized Column-Scheme

The limitations that can be over-seen are:

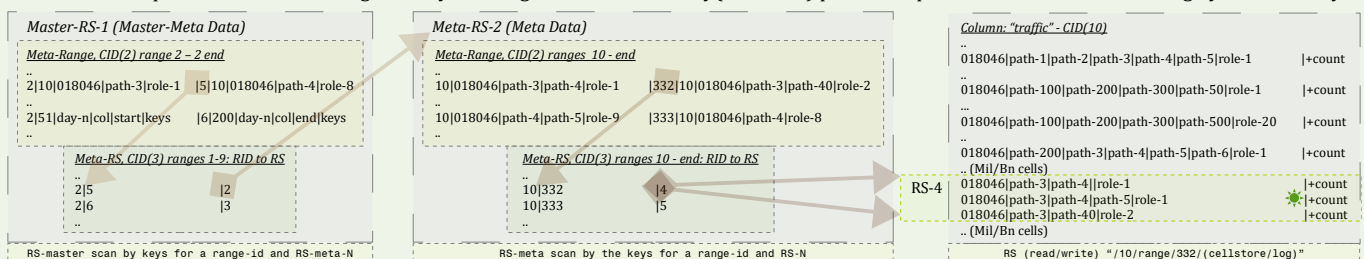
- ✗ Maximum number of columns, it is store-size of $\text{int64}(2^{64}) - 10$ (reserved cols) which can be improved by CID to be a string-type.
- ✗ Maximum size of Value or Key(after serialization), it is 4GB, while for such data size other limitations probably apply
- ✗ Maximum number of concurrent connections to a given server instance, it is the total available ports on the server, which can be further improved by using IPv6 and so as several IPs.

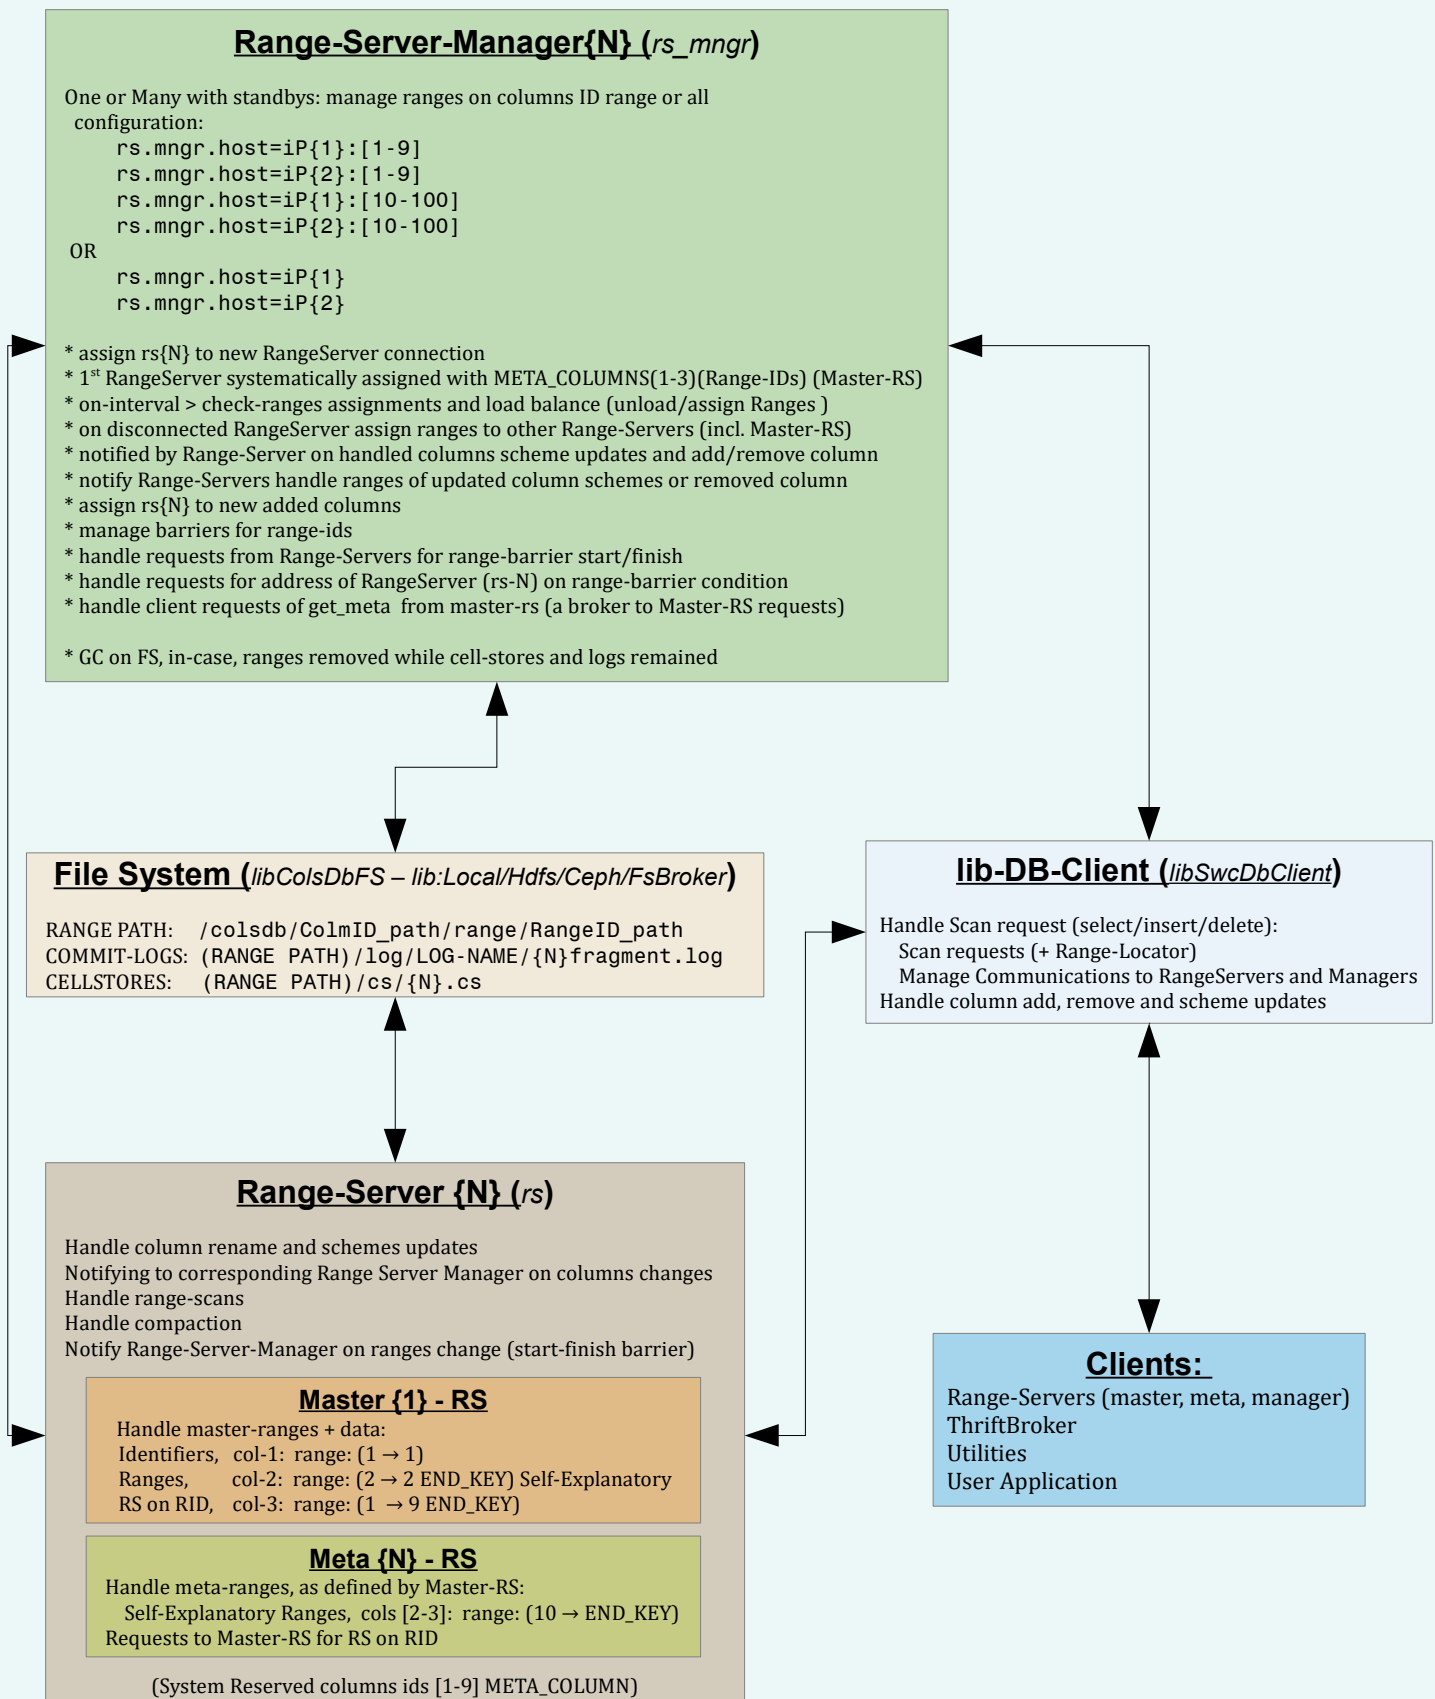
The capabilities to expect:

- ✓ Supposing one server can handle 2bil cells with one cell is being keys(1024B) and value(1024B) a 4TB in volume over 1300 ranges, to apply such base to RS-master means there are 500k RS-meta with 667mil RS-data with cell value being just a 4KB, that makes a SWC DB cluster in total handling 5.7ZB(zettabyte) of data volume and more on a compression ratio.
- ✓ A client can read at 100%(while Client's and RS's are equivalent) bandwidth, considering a perfect scan case of each client is requesting on different ranges, number of clients at a given time can be by the number of RS-data using 100% bandwidth each.

Some examples:

- Search indexing at <https://thither.direct/opensearch/> with Wide Column it is being row="sequences-of-words:domain:path" cf="lang" whereas with Super Wide Column it can be changed to keys=["sequences-of-words", "domain", "path", "lang"], makes the scan-select much optimized, especially if to query words-data of a domain & path, it would go on to ranges that start with domain & path skipping the seek through ranges of several other many domains that as well include the same word-sequences. While to have the same query on a Wide Column would require tripling the volume of data by using more indexes of word-sequences on a domain (and path) such as. row="domain:sequences:tripling" & row="domain:path:sequences". At current period the "open-search" on Thither.Direct does not offer querying data(words) on a site:domain or info:url-path as it is unreasonable over the data-volume overheads.
- A theoretical requirement for a building security tracking. Track of how many(an atomic-counter) personnel passed in an area of a building by role on a day:





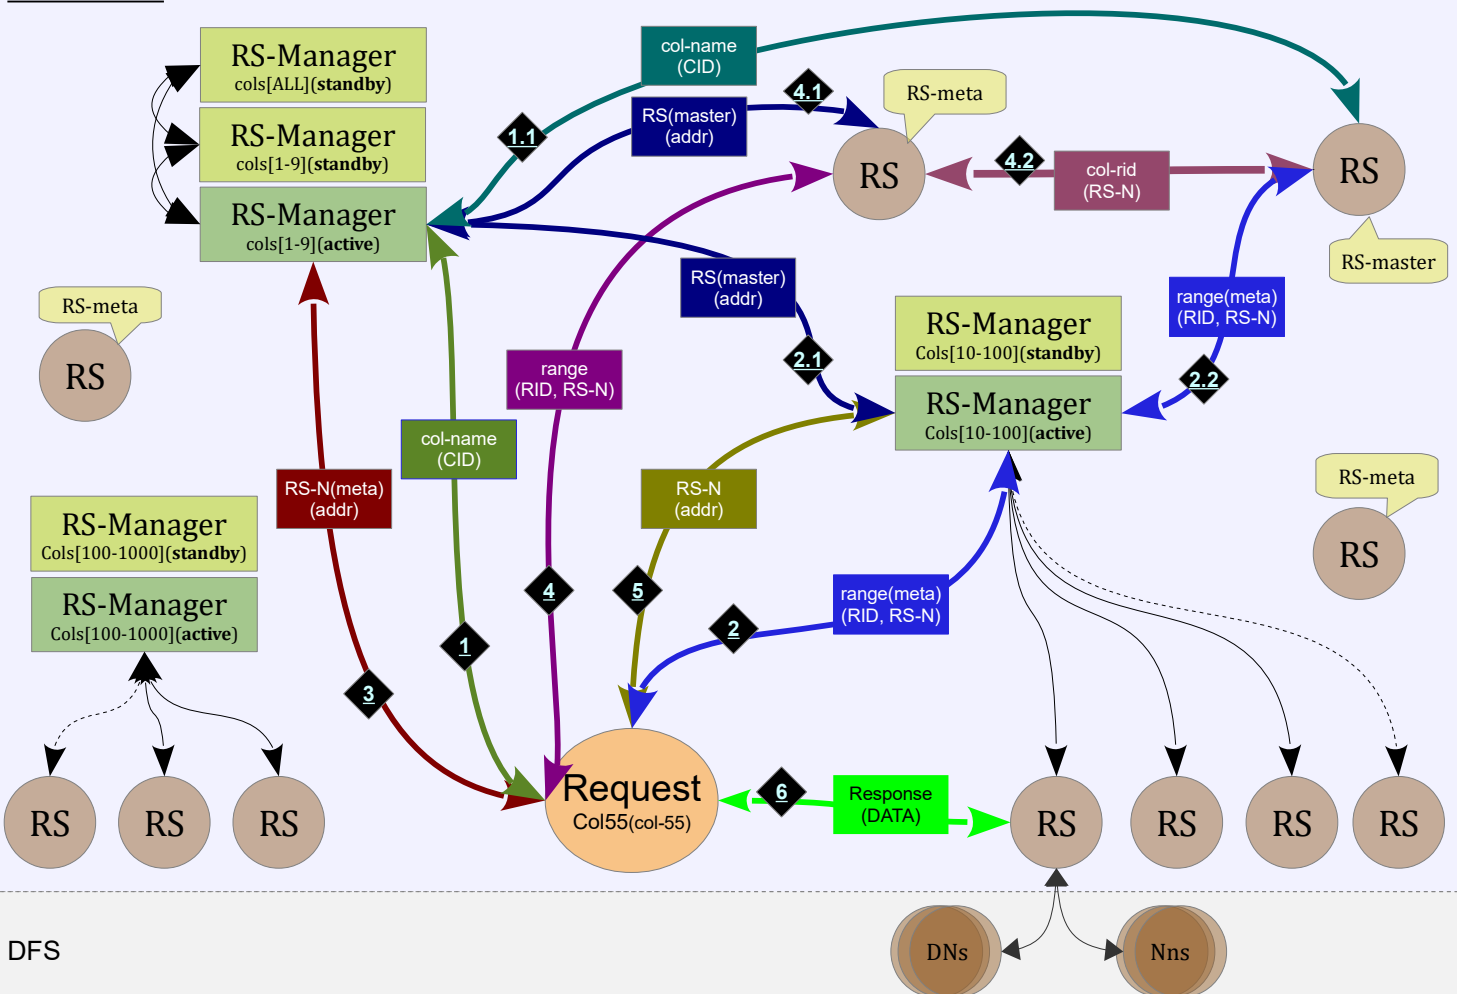
SWC DB: Failure Tolerance

author: Kashirin Alex
kashirin.alex@gmail.com

- ✓ A failed request to a RS-Manager is a connection fail-over to next in list from 'rs.mngr.host' configuration.
- ✓ A failed request to a RS(Master, Meta, Data)-N is fail-over to the new RS(addr) assign to RS-N by RS-Manager.
- ✓ RS-Manager, on interval or disconnection of a managed RS(either role), request to load ranges to another RS.
- ✓ RS-Master, as been a single instance, have connections only (in the case it is handling only master-ranges) with It's RS-Manager and requests for RS-N by column+rid from other RS-Managers{n} and RS{n}-Meta.
- ✓ Distribute File System, depends on the system and it's feature of routing to a datanode .
- ✓ RS(any) in case a connection or file-descriptor failure try to reconnect to the DFS.
- ✓ Communication over-heads of a resolved-data to column-name, RID-location or RS-address are kept while connection and data discard by TTL or a notification from RS-Managers on changes.

In worst case of outdated data being used with a request the RS return an error of a NOT_LOADED_RANGE.

DB-Client: *libColsDbClient.so*



SWC DB: LIB-DB-Client, Scan request (+ Range-Locator)

author: Kashirin Alex
kashirin.alex@gmail.com

Basic Process Flow of Scan request (+ Range-Locator)

Ranges Scan is done on per column base in-parallel(a client's max-range-locators config) with column's Scan Specifications

Scan-Specifications = cid, ScanSpecCellKeys(keys_start, keys_finish)

result = 0, last_cell_keys = NULL, new_start_keys = False, start_keys = ScanSpec.start_keys

DO scan_ranges_master:

```
1  get column-ID by name => (cid)
    RS-MANAGER[cid-2] - req, (cid-4, ["ReqColName"])
    get column-ID by name => (cid)
    RS-MASTER
    Scan-do, req (1-cell) => (cid)
2  get range-master-meta by (cid, start_keys, end_keys) => (rid, rs-N, next_meta_rid?)
    RS-MANAGER[cid]
    2.1 RS-MANAGER[cid-2]
        get RS-master-addr => (addr)
    2.2 RS-MASTER
        get range-master-meta => (rid, rs-N, next_meta_rid?)
        Scan-do (2-cell)(cid-2, ["cid", start_keys], ["cid", end_keys]) = rid(meta)
        Scan-do (1-cell)(cid-3, ["cid-2", "rid"]) = RS-N(meta)
    If no range-master-meta(rid, rs-N)
        goto finish
EXCEPT COMM:
    goto scan_ranges_master
```

DO scan_ranges_meta:

```
3  get RS-meta-addr by meta(rs-N, rid(barrier) => (addr)
    RS-MANAGER[cid-2] => (addr)
4  get range-meta by (cid, start_keys, end_keys) => (rid, rs-N, next_rid?)
    RS-META
    4.1 get range-meta by (cid, start_keys, end_keys) => (rid, rs-N)
        Scan-do (2-cell)(["cid", start_keys], ["cid", end_keys]) = rid
    4.2 Scan-request(cid-3, ["cid", "rid"]) = RS-N
    If no range-meta(rid, rs-N)
        if new_start_keys or next_meta_rid
            goto scan_ranges_master
        goto finish
EXCEPT COMM, NOT_LOADED_RANGE:
    goto scan_ranges_meta
```

DO scan_ranges_data:

```
5  get RS-addr by meta(rs-N, rid(barrier) => (addr)
    RS-MANAGER[cid-2] => (addr)
6  results = (RS-addr) scan-do(ScanSpecs)
    if no results
        if next_rid
            goto scan_ranges_meta
        goto finish
    (call_back) (available results), result+=results
    last_cell_keys=Result[-1]
EXCEPT COMM, NOT_LOADED_RANGE:
    goto scan_ranges_data
```

```
if result < limit(cell_limit) && last_cell_keys < start_keys:
    new_start_keys=True
    start_keys = last_cell->keys (all -ge on keys changed to -gt)
    goto scan_ranges_meta
DO finish:
    return result (call_back)
```

SWC DB: Query (SQL) scan

author: Kashirin Alex
kashirin.alex@gmail.com

```
select [where_clause [Columns-Intervals or Cells-Intervals]] [Flags(global-scope)];
```

Columns-Intervals: if not set, it is all columns from keys start to finish.

```
col(column-name-a1) = ( [Cells-Intervals] [and] [Cells-Intervals] [and] .. [Cells-Intervals] )
[and] ..
col(column-name-b1, .., column-name-b2) = ( [Cells-Intervals] [and] [Cells-Intervals] [and] .. [Cells-Intervals] )
```

Cells-Intervals: if not set, it is keys start to finish.

```
cells = ( [Cells-Interval] Flags(interval-scope) )
[and]
cells = ( [Cells-Interval] Flags(interval-scope) )
[and]..
cells = ( [Cells-Interval] Flags(interval-scope) )
```

Cells-Interval:

```
[ Condition-Keys ] [and] [ Condition-Value ] [and] [ Condition-Timestamp ]
```

Condition-Keys: keys comparator apply to every key that do not have a dedicated comparator, exact-match is keys=('k1', 'k2',,, 'kN')

```
keys [comparator] ( [comparator] "str-1", [comparator] "str-2", [comparator] "str-3", [comparator] "str-N" )
or (in-range)
([comparator] "str-1", [comparator] "str-N") [ <= or < ] keys [ <= or < ] ([comparator] "str-1", [comparator] "str-N")
```

Condition-Value:

```
value [comparator] "string"
or (for columns of counter type), not applicable comparators (prefix and regexp)
value [comparator] "int64_t(string)"
```

Condition-Timestamp: not applicable comparators (prefix and regexp)

```
timestamp [comparator] "YYYY/MM/DD HH:MM:ss.mmmuuunnn"
or (in-range)
"YYYY/MM/DD HH:MM:ss.mmmuuunnn" [ <= or < ] timestamp [ <= or < ] "YYYY/MM/DD HH:MM:ss.mmmuuunnn"
```

Comparator:

```
[ ^= ] : prefix (starts-with)
[ > ] : -gt (greater-than)
[ >= ] : -ge (greater-equal)
[ = ] : -eq (equal)
[ <= ] : -le (lower-equal)
[ < ] : -lt (lower-than)
[ != ] : -ne (not-equal)
[ re ] : regexp (regular-expression)

* -gt, -ge, -le, -lt are a bit-wise comparison
```

Flags: at global-scope apply to Cells-Interval flags to which does not have flags definitions

```
[ keys_only ] = TRUE on set # default FALSE
[ return_deletes ] = TRUE on set # default FALSE
[ limit ] = NUMBER(uint32_t) # default ALL
[ limit_by ] = "KEYS" or "." # default KEYS
[ offset ] = NUMBER(uint32_t) # default 0
[ offset_by ] = "KEYS" or "." # default KEYS
[ max_versions ] = NUMBER(uint32_t) # default ALL
```

An Example:

```
select
where
  col(ColNameA1) = (
    cells = ( (>='1-') <= keys = (<='1-1-', "1") and value = "Value-Data-1" and timestamp > "2010/05/29" limit=10 limit_by="KEYS" )
  )
and
  col(ColNameB1, ColNameB2) = (
    cells = ( (>='2-') <= keys = (<='2-2-', "1") and value = "Value-Data-2" and timestamp > "2010/05/29" )
    and
    cells = ( keys = (<='21-', "1") and timestamp > "2010/05/29" )
  )
max_versions=1;
```

SWC DB: Scan Specs & Results

author: Kashirin Alex
kashirin.alex@gmail.com

Scan Specs, lib-DB-Client:

ScanSpec (ListColumns columns; Flags flags;)	ColumnIntervals (int64_t oid ListCellsInterval cells_intervals) The object-type is applied to the range-locator (Client)	CellsInterval (Keys keys_start, keys_finish; Value value; Timestamp ts_start, ts_finish; Flags flags;)	Keys (ListKeys keys;)
Flags (uint32_t limit, offset, max_versions; LimitType limit_by, offset_by; bool return_deletes, keys_only;)	Key (const char* key; size_t key_len; Comparator comp;)	Value (const char* value; size_t value_len; Comparator comp;)	Timestamp (int64_t ts; Comparator comp;)

Scan Response, lib-DB-Client:

Result (List<Col> cols // ResponseFlag status = OK/PARTIAL/ERROR // Strings error_rs = ["N",])	Col (String name String id List<Cell> cells)	Cell (Strings keys int64_t timestamp char-array value uint32_t value_len)
---	---	---

SWC DB: Column Schema & Actions on Columns

author: Kashirin Alex
kashirin.alex@gmail.com

Although, there are schemas in the SWC-DB these can be considered as schema-less, exception to TTL, Counter and Max-Versions at the Cells level.

Location of Schema and Column-Name to ID:

Reserved Column-ID 4 with cell:

Keys = ("FLAG" "Col-Name", "CID")

Value = (value-serialized)version(int8)

|counter(int8)|ttl(int32)|compression(int8)|bloomfilter(int8)|max_versions(int32)|time_order(int8)|replication(int8)|blocksize(int32)

Configuration Options:

The following configurations available in the Column-Schema:

- CELL-LEVEL:
 - TTL: (int32_t) seconds
 - COUNTER: (bool), default False
 - MAX_VERSIONS: (int32_t), default 1 - not applicable with COUNTER
 - TIME_ORDER: ASC/DESC, default ASC – applied to order of MAX_VERSIONS
- BLOCK-LEVEL:
 - COMPRESSION: none/snappy/zlib/zstd/bmz/lzo/quicklz
 - BLOCKSIZE: (int32_t)
- CELLSSTORE-LEVEL:
 - REPLICATION: (int8_t) – replication factor applied to the DFS supporting file-replication , default 3
 - BLOOMFILTER: (int8_t) – none/all-keys

Adding a Column

SQL:

```
add column (  
    NAME="string",  
    COUNTER=bool,  
    MAX_VERSIONS=number,  
    TTL=number,  
    COMPRESSION="string",  
    BLOCKSIZE=number,  
    REPLICATION=number,  
    BLOOMFILTER="string"  
);
```

lib-DB-Client:

```
ColumnSpec (  
    String          name  
    int64_t         cid  
    Bool            counter  
    int32_t         max_versions  
    int32           ttl  
    Compressor::ENUM compression  
    int8_t          replication  
    BloomFilter     bloomfilter(TYPE::ENUM, factor, functions)  
)
```